



# Machine Learning

**Dr. Shahid Mahmood Awan**

**Assistant Professor**

**School of Systems and Technology, University of Management and Technology**

**shahid.awan@umt.edu.pk**

**Umer Saeed**(MS Data Science, BSc Telecommunication Engineering)

**Sr. RF Optimization & Planning Engineer**

**f2017313017@umt.edu.pk**



# Classification

Representation

# Classification


---

- ▶ **Email:** Spam/Not Spam?
- ▶ **Online Transactions:** Fraudulent (Yes/No)?
- ▶ **Tumor:** Malignant/Benign?

$$y \in \{0,1\}$$


- ▶ **0:** "Negative Class" (e.g. benign tumor)
- ▶ **1:** "Positive Class" (e.g. malignant tumor)
- ▶ and they are sometimes also denoted by the symbols "-" and "+".
- ▶ Such problem also know as "Binary Classification".
- ▶ Predict value  $y$  is discrete in classification.

# Classification



- ▶ Threshold classifier output  $h_{\theta}(x)$  at 0.5
  - If  $h_{\theta}(x) \geq 0.5$ , predict " $y = 1$ "
  - If  $h_{\theta}(x) < 0.5$ , predict " $y = 0$ "

# Classification



- ▶ To attempt classification, one method is to use linear regression and map all predictions greater than 0.5 as a 1 and all less than 0.5 as a 0. However, this method doesn't work well because classification is not actually a linear function.

# Classification

---

*Classification:  $y = 0$  or  $1$*

*$h_{\theta}(x)$  can be  $> 1$  or  $< 0$*

*Logistic Regression:  $0 \leq h_{\theta}(x) \leq 1$*

- ▶ Logistic Regression is also known as Classification Algorithm.

# Question

---

- ▶ Which of the following statements is true?
  
- ▶ (a) If linear regression doesn't work on a classification task as in the previous example shown in the video, applying feature scaling may help.
  
- ▶ (b) If the training set satisfies  $0 \leq y^{(i)} \leq 1$  for every training example  $(x^{(i)}, y^{(i)})$ , then linear regression's prediction will also satisfy  $0 \leq h(x) \leq 1$  for all values of  $x$ .
  
- ▶ (c) If there is a feature  $x$  that perfectly predicts  $y$ , i.e. if  $y=1$  when  $x \geq c$  and  $y=0$  whenever  $x < c$  (for some constant  $c$ ), then linear regression will obtain zero classification error.
  
- ▶ **(d) None of the above statements are true. (Answer)**



# **Classification**

## **Hypothesis Representation**




# Hypothesis Representation

**Want** : Logistic Regression:  $0 \leq h_{\theta}(x) \leq 1$

$$h_{\theta}(x) = \mathbf{g}(\theta^T \mathbf{x}) \quad g(z) = \frac{1}{1 + e^{-z}} ; \text{Logistic/Sigmoid function}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T \mathbf{x}}}$$

- The function  $g(z)$ , shown here, maps any real number to the  $(0, 1)$  interval, making it useful for transforming an arbitrary-valued function into a function better suited for classification.



# Hypothesis Representation

---

► **Remember.**

$$\begin{aligned} z = 0, e^0 = 1 &\Rightarrow g(z) = 1/2 \\ z \rightarrow \infty, e^{-\infty} \rightarrow 0 &\Rightarrow g(z) = 1 \\ z \rightarrow -\infty, e^{\infty} \rightarrow \infty &\Rightarrow g(z) = 0 \end{aligned}$$

# Hypothesis Representation

---

- ▶ **Interpretation of Hypothesis Output**

$h_{\theta}(x) = \text{estimated probability that } y = 1 \text{ on input } x$

- ▶ **Example**

$$\text{If } x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ \text{tumorSize} \end{bmatrix}$$

- ▶ Tell patient that 70% chance of tumor being malignant.

$$P(y = 0|x; \theta) + P(y = 1|x; \theta) = 1$$

$$P(y = 1|x; \theta) = 1 - P(y = 0|x; \theta)$$

# Question

---

- ▶ Suppose we want to predict, from data  $x$  about a tumor, whether it is malignant ( $y=1$ ) or benign ( $y=0$ ). Our logistic regression classifier outputs, for a specific tumor,  $h(x)=P(y=1|x;\theta)=0.7$ , so we estimate that there is a 70% chance of this tumor being malignant. What should be our estimate for  $P(y=0|x;\theta)$ , the probability the tumor is benign?

$(a)P(y = 0|x; \theta) = 0.3$  Answer

$(b)P(y = 0|x; \theta) = 0.7$

$(c)P(y = 0|x; \theta) = 0.7^2$

$(d)P(y = 0|x; \theta) = 0.3 \times 0.7$



**Classification**

**Decision Boundary**

# Decision Boundary

$$h_{\theta}(x) = \mathbf{g}(\theta^T x) \quad g(z) = \frac{1}{1 + e^{-z}}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

**Suppose predict "y=1" if  $h_{\theta}(x) \geq 0.5$**


$$g(z) \geq 0.5 \text{ when } z \geq 0$$

$$h_{\theta}(x) = g(\theta^T x) \geq 0.5 \text{ whenever } \theta^T x \geq 0$$


**Predict "y=0" if  $h_{\theta}(x) < 0.5$**

$$g(z) < 0.5 \text{ when } z < 0$$

$$h_{\theta}(x) = g(\theta^T x) < 0.5 \text{ whenever } \theta^T x < 0$$



# Decision Boundary



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

$$\text{Let } \theta_0 = -3, \theta_1 = 1, \theta_2 = 1$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$$

$$\text{Predict } y=1 \text{ if } -3 + x_1 + x_2 \geq 0$$

$$x_1 + x_2 \geq 3$$


$$y=0 \text{ if } x_1 + x_2 < 3$$

- The **decision boundary** is the line that separates the area where  $y = 0$  and where  $y = 1$ . It is created by our hypothesis function.

# Non-Linear Decision Boundaries

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

$$\text{Let } \theta_0 = -1, \theta_1 = 0, \theta_2 = 0, \theta_3 = 1, \theta_4 = 1$$



$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \theta_3 \\ \theta_4 \end{bmatrix} = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

$$\text{Predict } y=1 \text{ if } -1 + x_1^2 + x_2^2 \geq 0$$

$$x_1^2 + x_2^2 \geq 1$$

$$y = 0 \text{ if } x_1^2 + x_2^2 < 1$$

- Decision Boundary is the property of the hypothesis and parameter not a training set.





# **Logistic Regression Model**

## **Cost Function**

# Logistic Regression Cost Function

---

*Training Set:  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots (x^{(m)}, y^{(m)})\}$*

$$m \text{ examples: } x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ \vdots \\ x_n \end{bmatrix}$$

$$\text{Hypothesis: } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

*How to choose parameters  $\theta$ ?*

# Logistic Regression Cost Function

---

► **Linear regression:**

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^i) - y^{(i)})^2$$


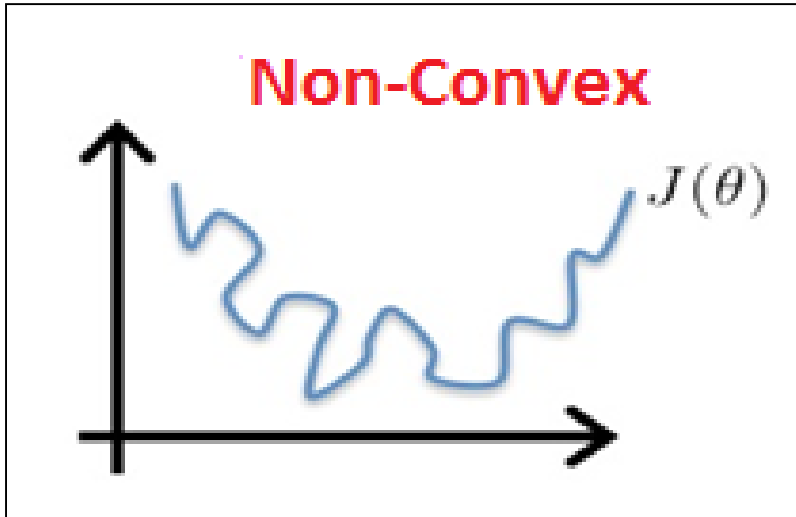
$$\text{Cost}(h_{\theta}(x^i), y^{(i)}) = \frac{1}{2} (h_{\theta}(x^i) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^i), y^{(i)})$$

$$\text{Logistic Regression Hypothesis: } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$



# Logistic Regression Cost Function

---



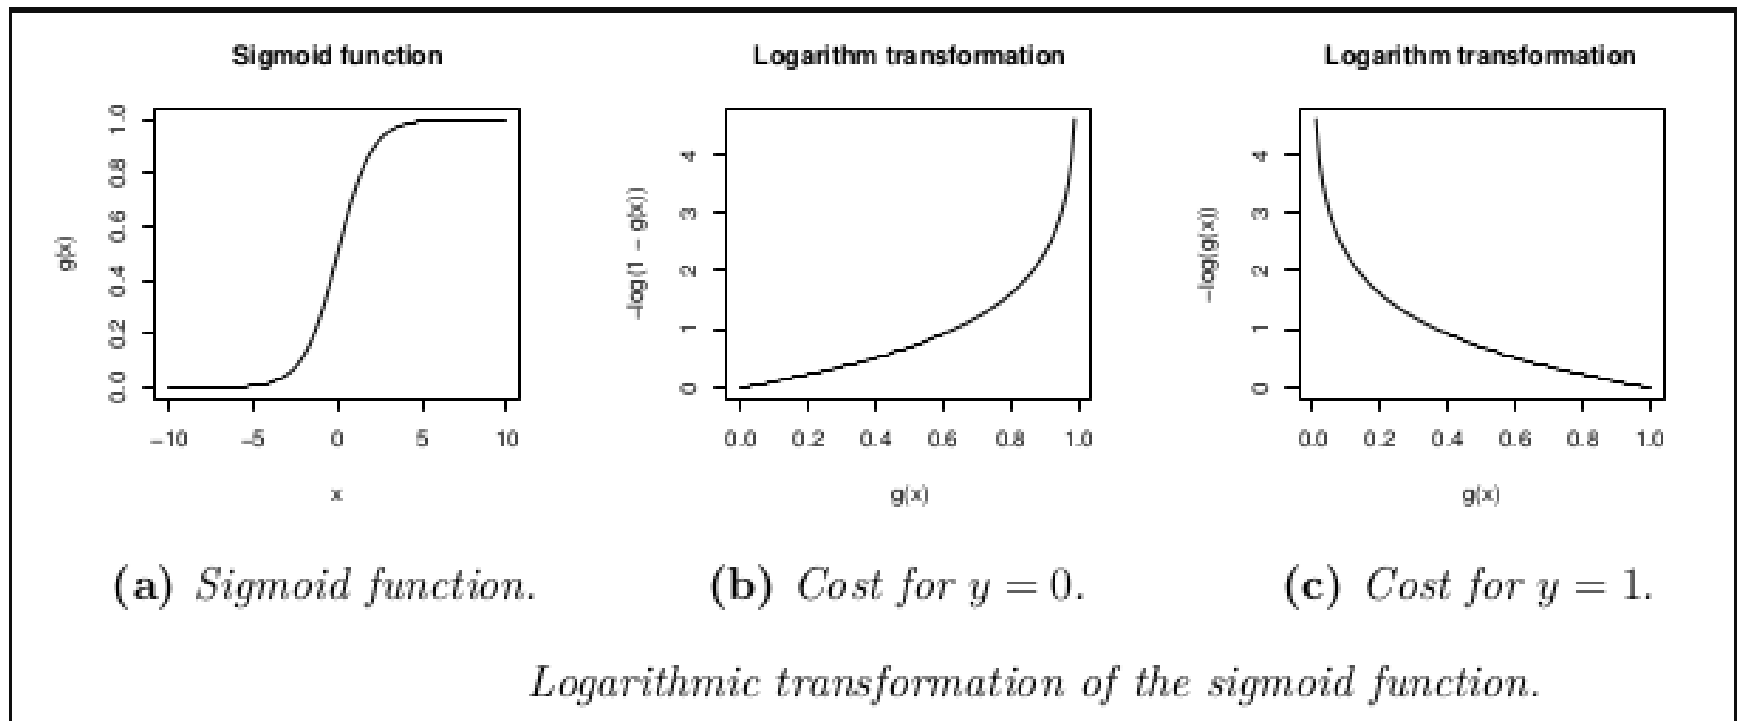
# Logistic Regression Cost Function

- ▶ **Question**
- ▶ Consider minimizing a cost function  $J(\theta)$ . Which one of these functions is convex?



# Logistic Regression Cost Function

$$Cost(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$




# Logistic Regression Cost Function

$Cost(h_{\theta}(x), y) = -\log(h_{\theta}(x))$  if  $y = 1$

- ▶ **Case-1:** If  $(h_{\theta}(x))=0$   $Cost(h_{\theta}(x), y) = -\log(0) = \infty$
- ▶ **Case-2:** If  $(h_{\theta}(x))=1$   $Cost(h_{\theta}(x), y) = -\log(1) = 0$
- ▶ If our correct answer 'y' is 1, then the cost function will be 0 if our hypothesis function outputs 1. If our hypothesis approaches 0, then the cost function will approach infinity.
- ▶ Captures intuition that if  $(h_{\theta}(x))=0$ , (predict  $P(y=1|x;\theta)$ , but  $y=1$  we'll penalize learning algorithm by a very large cost.


When  $y = 1$ , we get the following plot for  $J(\theta)$  vs  $h_{\theta}(x)$ :



# Logistic Regression Cost Function

$Cost(h_{\theta}(x), y) = -\log(1 - h_{\theta}(x))$  if  $y = 0$

- ▶ **Case-1:** If  $(h_{\theta}(x))=0$       $Cost(h_{\theta}(x), y) = -\log(1 - 0) = -\log(1) = 0$
- ▶ **Case-2:** If  $(h_{\theta}(x))=1$       $Cost(h_{\theta}(x), y) = -\log(1 - 1) = -\log(0) = \infty$
- ▶ If our correct answer 'y' is 0, then the cost function will be 0 if our hypothesis function also outputs 0. If our hypothesis approaches 1, then the cost function will approach infinity.





# Question

- In logistic regression, the cost function for our hypothesis outputting (predicting)  $h(x)$  on a training example that has label  $y \in \{0, 1\}$  is;

☒ If  $h_{\theta}(x) = y$ , then  $\text{cost}(h_{\theta}(x), y) = 0$  (for  $y = 0$  and  $y = 1$ ).

Correct

☒ If  $y = 0$ , then  $\text{cost}(h_{\theta}(x), y) \rightarrow \infty$  as  $h_{\theta}(x) \rightarrow 1$ .

Correct

☐ If  $y = 0$ , then  $\text{cost}(h_{\theta}(x), y) \rightarrow \infty$  as  $h_{\theta}(x) \rightarrow 0$ .

Un-selected is correct

☒ Regardless of whether  $y = 0$  or  $y = 1$ , if  $h_{\theta}(x) = 0.5$ , then  $\text{cost}(h_{\theta}(x), y) > 0$ .

Correct



# **Logistic Regression Model**

**Simplified Cost Function and gradient descent**

# Logistic Regression Cost Function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^i), y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

- Note:  $y=0$  or  $y=1$  always.

$$\text{Cost}(h_{\theta}(x), y) = -y(\log(h_{\theta}(x))) - (1 - y) \log(1 - h_{\theta}(x))$$

$$\text{Cost}(h_{\theta}(x), y) = -[y(\log(h_{\theta}(x))) + (1 - y) \log(1 - h_{\theta}(x))]$$

**Case – 1 ( $y = 0$ )**

$$\text{Cost}(h_{\theta}(x), y) = -(0)(\log(h_{\theta}(x))) - (1 - 0) \log(1 - h_{\theta}(x)) = -\log(1 - h_{\theta}(x))$$

**Case – 1 ( $y = 1$ )**

$$\text{Cost}(h_{\theta}(x), y) = -(1)(\log(h_{\theta}(x))) - (1 - 1) \log(1 - h_{\theta}(x)) = -\log(h_{\theta}(x))$$

# Logistic Regression Cost Function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^i), y^{(i)})$$

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y(\log(h_{\theta}(x))) + (1 - y) \log(1 - h_{\theta}(x))]$$

- ▶ This cost function is derived from statistics using the principle of maximum likelihood estimation.
- ▶ The property of maximum likelihood estimation that is “convex”.

**To fit parameters  $\theta$ :**  $\min_{\theta} (J(\theta))$

- ▶ To make a prediction given new  $x$ ;

$$\text{Logistic Regression Hypothesis: } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

$h_{\theta}(x)$  = estimated probability that  $y = 1$  on input  $x$  ( $P(y = 1|x; \theta)$ )

# Logistic Regression Cost Function

---

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y(\log(h_{\theta}(x)) + (1 - y) \log(1 - h_{\theta}(x))) \right]$$

- ▶ We want  $\min_{\theta} (J(\theta))$ .
- ▶ To do so, we use gradient descent, but we first need to find the partial derivatives  $\frac{\partial}{\partial \theta_j}$

# Logistic Regression Cost Function

- ▶ We're going to make use of a neat property of the logistic functions:

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$\frac{d}{dx} \left( \frac{u}{v} \right) = \frac{v \frac{du}{dx} - u \frac{dv}{dx}}{v^2}$$

$$\frac{d}{dx} e^x = e^x$$

- ▶ By taking the derivatives on the both sides, we have

$$h_{\theta}(\acute{x}) = \frac{d}{dx} \left[ \frac{1}{1 + e^{-\theta^T x}} \right]$$

$$h_{\theta}(\acute{x}) = \frac{\left(1 + e^{-\theta^T x}\right) \frac{d}{dx} (1) - (1) \frac{d}{dx} \left(1 + e^{-\theta^T x}\right)}{\left(1 + e^{-\theta^T x}\right)^2}$$

$$h_{\theta}(\acute{x}) = \frac{\left(1 + e^{-\theta^T x}\right) (0) - \left[ \frac{d}{dx} (1) - \frac{d}{dx} (e^{-\theta^T x}) \right]}{\left(1 + e^{-\theta^T x}\right)^2}$$

# Logistic Regression Cost Function

---

$$h_{\theta}(\acute{x}) = \frac{e^{-\theta^T x}}{(1 + e^{-\theta^T x})^2}$$

$$h_{\theta}(\acute{x}) = \frac{1 + e^{-\theta^T x} - 1}{(1 + e^{-\theta^T x})^2}$$

$$h_{\theta}(\acute{x}) = \frac{1}{1 + e^{-\theta^T x}} - \frac{1}{(1 + e^{-\theta^T x})^2}$$

$$h_{\theta}(\acute{x}) = \frac{1}{1 + e^{-\theta^T x}} \left[ 1 - \frac{1}{1 + e^{-\theta^T x}} \right]$$

$$h_{\theta}(\acute{x}) = h_{\theta}(x)[1 - h_{\theta}(x)]$$

# Logistic Regression Cost Function

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y(\log(h_{\theta}(x)) + (1 - y) \log(1 - h_{\theta}(x))) \right]$$

- By taking the derivatives on the both sides, we have

$$\begin{aligned} \frac{\partial}{\partial \theta_j} J(\theta) &= -\frac{1}{m} \left[ \frac{\partial}{\partial \theta_j} \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x)) \right] \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \frac{1}{h_{\theta}(x^{(i)})} \frac{\partial}{\partial \theta_j} h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \frac{1}{1 - h_{\theta}(x^{(i)})} \left( -\frac{\partial}{\partial \theta_j} h_{\theta}(x^{(i)}) \right) \right] \\ \frac{\partial}{\partial \theta_j} J(\theta) &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \frac{x_j^{(i)}}{h_{\theta}(x^{(i)})} h_{\theta}(x^{(i)})(1 - h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \frac{x_j^{(i)}}{1 - h_{\theta}(x^{(i)})} h_{\theta}(x)(1 - h_{\theta}(x^{(i)})) \right] \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} x_j^{(i)} (1 - h_{\theta}(x^{(i)})) - (1 - y^{(i)}) x_j^{(i)} h_{\theta}(x^{(i)}) \right] \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} x_j^{(i)} - x_j h_{\theta}(x^{(i)}) \right] \\ \frac{\partial}{\partial \theta_j} J(\theta) &= \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \end{aligned}$$



# Gradient Descent

---

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial_j} (J(\theta))$$

(Simultaneously update all  $\theta_j$ )

}

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^i) x_j^{(i)})$$

(Simultaneously update all  $\theta_j$ )

}

- ▶ Algorithm looks identical to linear regression!



# **Logistic Regression Model**

## **Advanced Optimization**

# Advanced Optimization

---

- ▶ 1- "Conjugate gradient"
- ▶ 2- "BFGS"
- ▶ 3- "L-BFGS"
- ▶ are more sophisticated, faster ways to optimize  $\theta$  that can be used instead of gradient descent.



# **Logistic Regression Model**

**Multi-Class Classification: One-Vs-all**


# Multiclass Classification: One-vs-all

- ▶ Now we will approach the classification of data when we have more than two categories. Instead of  $y = \{0,1\}$  we will expand our definition so that  $y = \{0,1,...n\}$ .
- ▶ **Examples:**


Code(format-1)	Code(format-2)	Email Foldering/Tagging	Medical Diagrams	Weather
0	1	Work	Not Till	Sunny
1	2	Friends	Cold	Cloudy
2	3	Family	Flu	Rain
3	4	Hobby		Snow

# Multiclass Classification: One-vs-all


Binary classification:




Multi-class classification:



# Multiclass Classification: One-vs-all (one-vs-Rest)




$h_{\theta}^{(1)}$ : Superscript 1 is stands for class – "1"




$$P(y = 1|x; \theta)$$

# Multiclass Classification: One-vs-all (one-vs-Rest)




$h_{\theta}^{(2)}$ : Superscript 2 stands for class – "2"




$$P(y = 2|x; \theta)$$



# Multiclass Classification: One-vs-all (one-vs-Rest)



$h_{\theta}^{(3)}$ : Superscript 3 is stands for class – "3"



$$P(y = 3|x; \theta)$$

$$h_{\theta}^{(i)}(x) = P(y = i|x; \theta) \quad (i = 1, 2, 3)$$

# Multiclass Classification: One-vs-all (one-vs-Rest)

---



- ▶ Train a logistic regression classifier  $h_{\theta}^{(i)}(x)$  for each class  $i$  to predict the probability that  $y = i$
- ▶ On a new input  $x$ , to make a prediction, pick the class  $i$  that maximizes.

$$\max_i h_{\theta}^{(i)}(x)$$

# Multiclass Classification: One-vs-all

---

- ▶ Since  $y = \{0, 1 \dots n\}$ , we divide our problem into  $n+1$  (+1 because the index starts at 0) binary classification problems; in each one, we predict the probability that 'y' is a member of one of our classes.

$$\begin{aligned} y &\in \{0, 1 \dots n\} \\ h_{\theta}^{(0)}(x) &= P(y = 0|x; \theta) \\ h_{\theta}^{(1)}(x) &= P(y = 1|x; \theta) \\ &\dots \\ h_{\theta}^{(n)}(x) &= P(y = n|x; \theta) \\ \text{prediction} &= \max_i (h_{\theta}^{(i)}(x)) \end{aligned}$$

- ▶ We are basically choosing one class and then lumping all the others into a single second class.
- ▶ We do this repeatedly, applying binary logistic regression to each case, and then use the hypothesis that returned the highest value as our prediction.

# Question

---

- ▶ Suppose you have a multi-class classification problem with  $k$  classes (so  $y \in \{1, 2, \dots, k\}$ ). Using the 1-vs.-all method, how many different logistic regression classifiers will you end up training?

(a)  $k - 1$

(b)  $k$  (answer)

(c)  $k + 1$

(d) Approximately  $\log_2(k)$



# **Logistic Regression Model**

## **The Problem of Over fitting**


# The Problem of Over fitting (Linear Regression)

- ▶ Consider the problem of predicting  $y$  from  $x \in \mathbb{R}$ .

- ▶ The figure shows the result of fitting a  $y = \theta_0 + \theta_1 x$  to a dataset.


- ▶ We see that the data doesn't really lie on straight line, and so the fit is not very good.

- ▶ we'll say the figure shows an instance of **under-fitting**—in which the data clearly shows structure not captured by the model.



# The Problem of Over fitting (Linear Regression)


- ▶ Instead, if we had added an extra feature  $x^2$ , and fit  $y = \theta_0 + \theta_1 x + \theta_2 x^2$ , then we obtain a slightly better fit to the data.
- ▶ Naively, it might seem that the more features we add, the better.



# The Problem of Over fitting (Linear Regression)

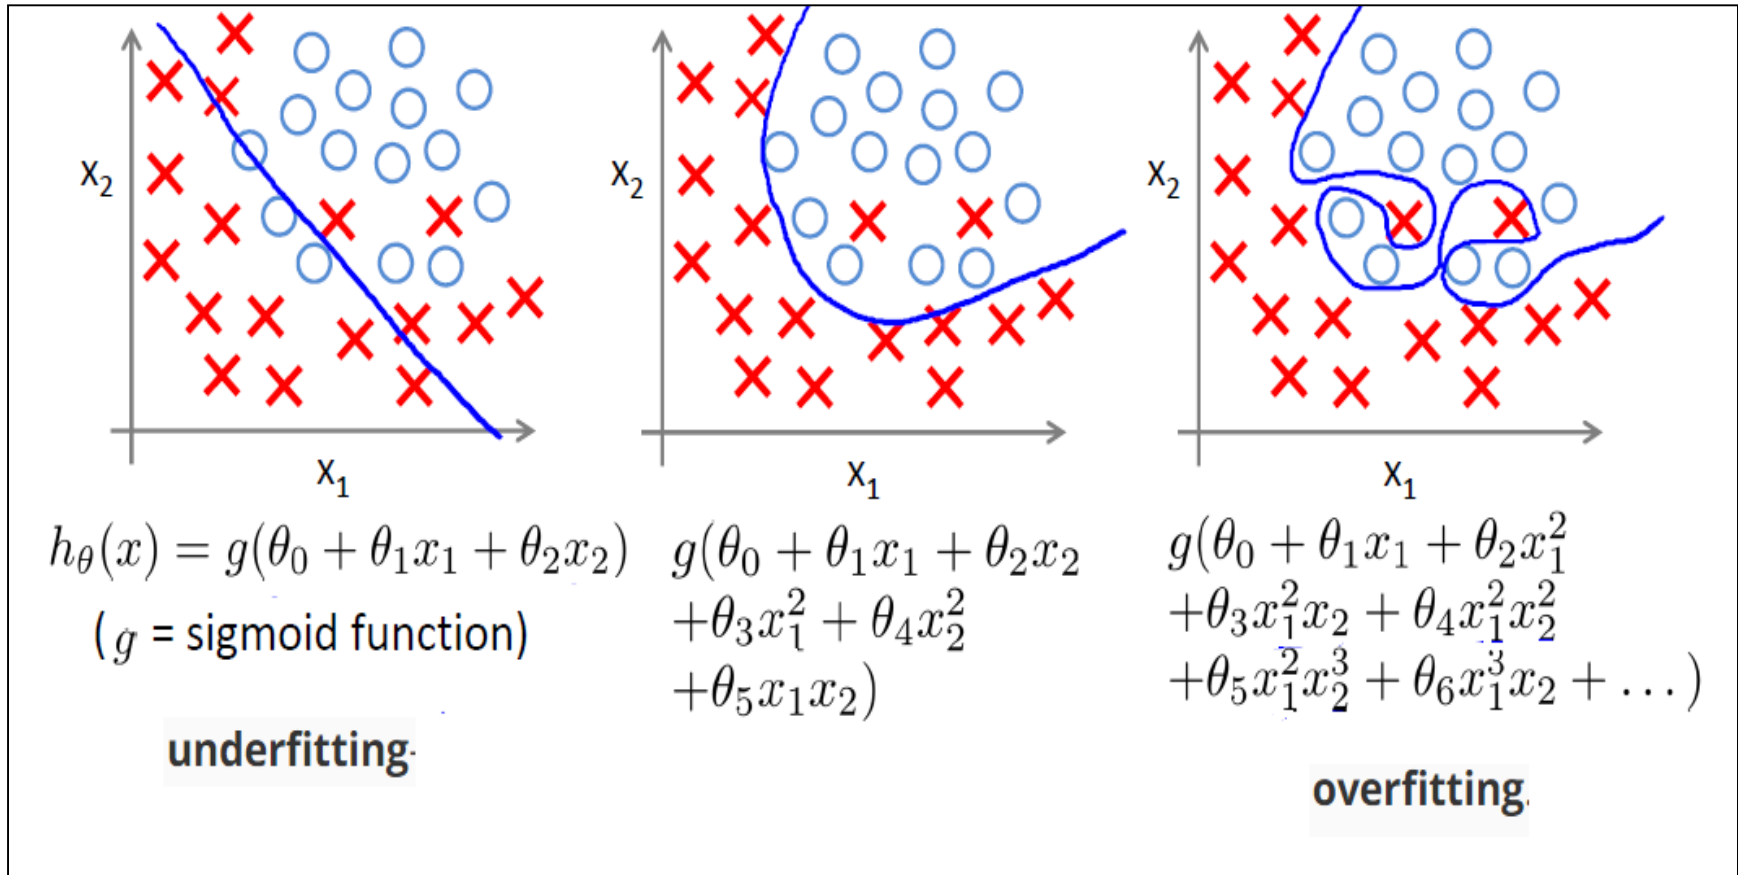


- ▶ There is also a danger in adding too many features: The figure is the result of fitting a 5<sup>th</sup> order polynomial
- ▶ We see that even though the fitted curve passes through the data perfectly, we would not expect to be a very good predictor.
- ▶ we'll say the figure shows an instance of **over-fitting**—in which the data clearly shows structure not captured by the model.





# The Problem of Over fitting (Logistic Regression)



- ▶ This terminology is applied to both linear and logistic regression

# The Problem of Over fitting

---

- ▶ **Under-fitting, or high bias**, is when the form of our hypothesis function  $h$  maps poorly to the trend of the data.
- ▶ It is usually caused by a function that is too simple or uses too few features.
- ▶ At the other extreme, **over-fitting, or high variance**, is caused by a hypothesis function that fits the available data but does not generalize well to predict new data.

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h(x)^{(i)} - y^{(i)})^2 \approx 0$$

- ▶ It is usually caused by a complicated function that creates a lot of unnecessary curves and angles unrelated to the data.

# Addressing Over-fitting

---

- ▶ There are two main options to address the issue of over-fitting:
- ▶ **1- Reduce number of features.**
  - ▶ - Manually select which features to keep.
  - ▶ - *Model selection algorithm*
- ▶ **2- Regularization**
  - ▶ - Keep all the features, but reduce the magnitude/values of parameters  $\theta_j$
  - ▶ - Works well when we have a lot of features, each of which contributes a bit to predicting  $y$ .

# Question

---


- ▶ Consider the medical diagnosis problem of classifying tumors as malignant or benign. If a hypothesis  $h(x)$  has overfit the training set, it means that:
  - ▶ (a) It makes accurate predictions for examples in the training set and generalizes well to make accurate predictions on new, previously unseen examples.
  - ▶ (b) It does not make accurate predictions for examples in the training set, but it does generalize well to make accurate predictions on new, previously unseen examples.
  - ▶ (c ) It makes accurate predictions for examples in the training set, but it does not generalize well to make accurate predictions on new, previously unseen examples. (Answer)
  - ▶ (d) It does not make accurate predictions for examples in the training set and does not generalize well to make accurate predictions on new, previously unseen examples.




## **Linear Regression Model**

## **Regularization-Cost Function**

# Regularization



$$\theta_0 + \theta_1 x + \theta_2 x^2$$



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \cancel{\theta_3 x^3} + \cancel{\theta_4 x^4}$$

Suppose we penalize and make  $\theta_3, \theta_4$  really small.

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + 1000 \theta_3^2 + 1000 \theta_4^2$$

$$\theta_3 \approx 0$$

$$\theta_4 \approx 0$$

# Regularization

---

- ▶ Small values of parameters  $\theta_0, \theta_1, \dots, \theta_n$
- ▶ - “Simples” hypothesis
- ▶ - Less prone to over-fitting
  
- ▶ **Example (Housing)**

*Features:  $x_1, x_2, \dots, x_{100}$*

*Parameters:  $\theta_0, \theta_1, \dots, \theta_{100}$*

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

# Regularization

---

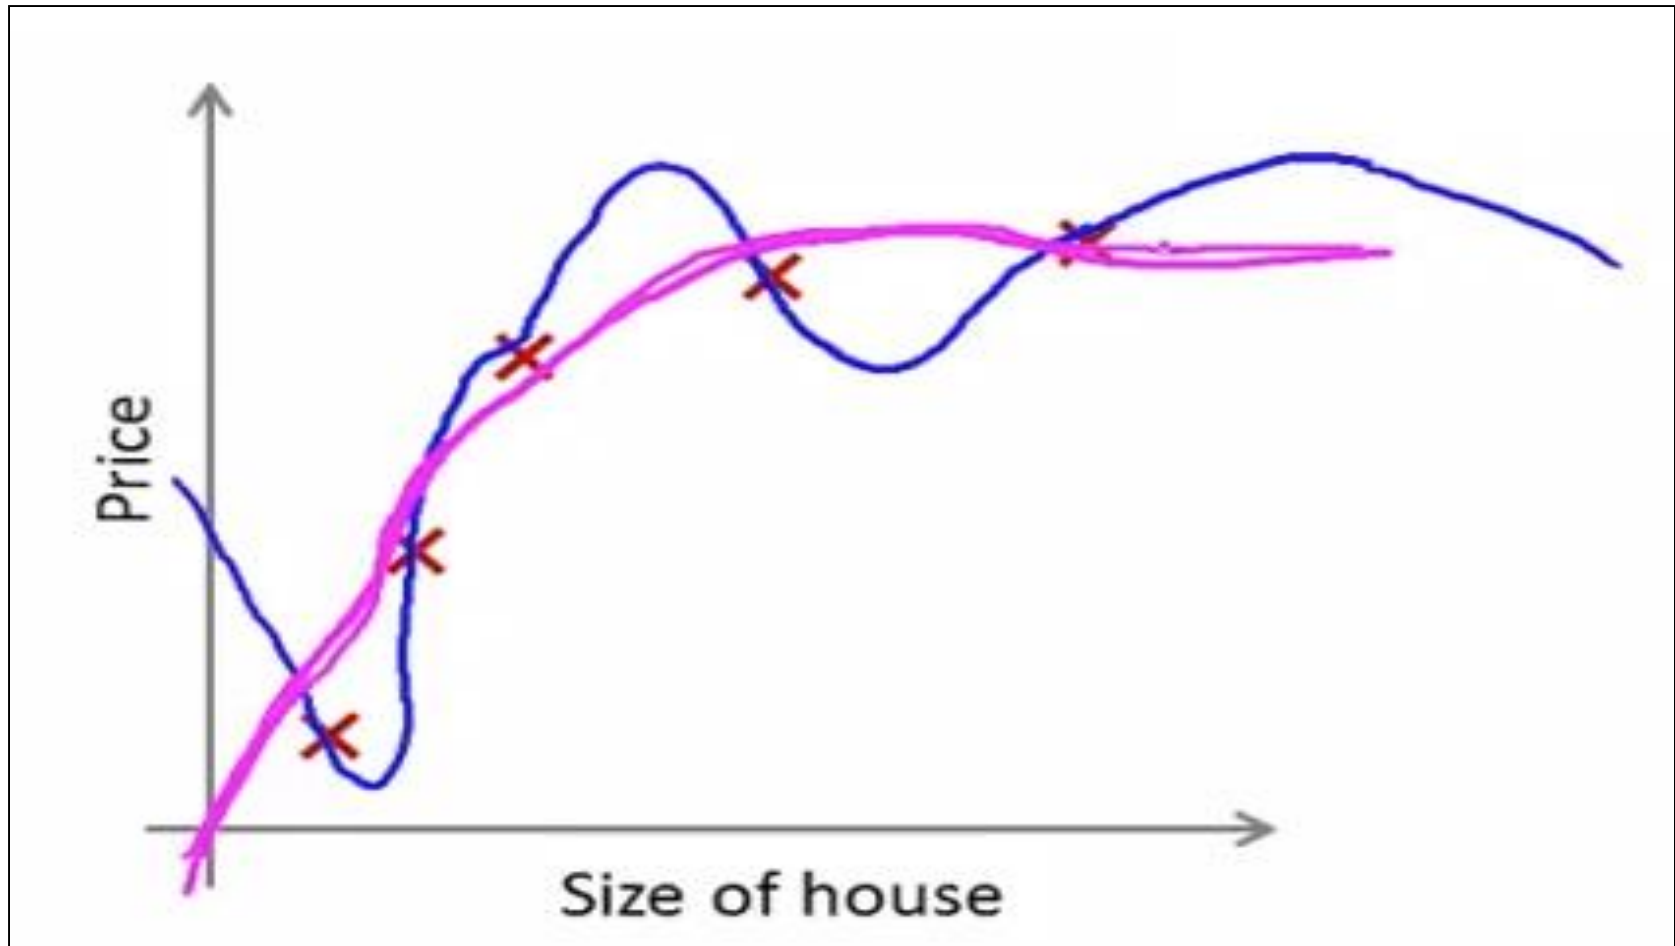
$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

- ▶ The  $\lambda$ , or lambda, is the **regularization parameter**. It determines how much the costs of our theta parameters are inflated.
- ▶  **$\lambda$  is control the trade off between 2 different goals:**
  - ▶ 1-  $\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$  fit the training data set well.
  - ▶ 2-  $\sum_{j=1}^n \theta_j^2$  Keeps the parameter small.
- ▶ and therefore keeping hypothesis relatively simple to avoid over-fitting.



# Regularization

## ► Example




# Regularization

- ▶ In regularization linear regression, we choose  $\theta$  to minimize;

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

- ▶ What if  $\lambda$  is set to an extremely large value (perhaps far too large for our problem, say  $\lambda = 10^{10}$ )?



# Question

In regularized linear regression, we choose  $\theta$  to minimize:

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

What if  $\lambda$  is set to an extremely large value (perhaps too large for our problem, say  $\lambda = 10^{10}$ )?

- ☐ Algorithm works fine; setting  $\lambda$  to be very large can't hurt it.
- ☐ Algorithm fails to eliminate overfitting.
- ☒ Algorithm results in underfitting (fails to fit even the training set).

Correct

- ☐ Gradient descent will fail to converge.

# Regularization

*Repeat*

$$\theta_0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0) \quad for (j = 0)$$

$$\left. \begin{aligned} &\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_j) \\ & \end{aligned} \right\} \text{----- } (\textbf{Equation - A}) \quad \text{for } (j = 1, 2, 3 \dots, n)$$

# Regularization

Calculation for  $\frac{\partial}{\partial \theta_0} J(\theta_0)$

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{\partial}{\partial \theta_0} \left[ \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{1}{2m} \frac{\partial}{\partial \theta_0} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

# Regularization

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{1}{2m} \left[ \frac{\partial}{\partial \theta_0} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \right] + \frac{\partial}{\partial \theta_0} \left[ \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{1}{2m} \left[ \frac{\partial}{\partial \theta_0} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2 \right] + \frac{\partial}{\partial \theta_0} \left[ \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{2}{2m} \left[ \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) \frac{\partial}{\partial \theta_0} (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) \right] + \frac{\partial}{\partial \theta_0} \left[ \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

# Regularization

---

$$\begin{aligned} & \frac{\partial}{\partial \theta_0} J(\theta) \\ &= \frac{1}{m} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) \left\{ \frac{\partial}{\partial \theta_0} (\theta_0) + \frac{\partial}{\partial \theta_0} (\theta_1 x^{(i)}) - \frac{\partial}{\partial \theta_0} (y^{(i)}) \right\} \right] \end{aligned}$$

# Regularization

---

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{1}{m} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) \{1\} \right]$$

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{1}{m} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) \{x_0^{(i)}\} \right]$$



# Regularization

Calculation for  $\frac{\partial}{\partial \theta_j} J(\theta_j)$

$$J(\theta) = \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \left[ \frac{1}{2m} \left[ \sum_{i=1}^m (h_{\theta}(x^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \left[ \frac{1}{2m} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

# Regularization

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{2m} \left[ \frac{\partial}{\partial \theta_j} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{2m} \left[ \frac{\partial}{\partial \theta_j} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2 \right] + \frac{\partial}{\partial \theta_j} \left[ \lambda \sum_{j=1}^n \theta_j^2 \right] \right]$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{2}{2m} \left[ \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) [x_j^{(i)}] \right] + \lambda \theta_j \right]$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) [x_j^{(i)}] + \lambda \theta_j \right]$$

# Regularization

---

By putting the values in equation (A), we have

Repeat

{

$$\theta_0 := \theta_0 - \frac{\alpha}{m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) \{x_0^{(i)}\}$$

$$\theta_j := \theta_j - \frac{\alpha}{m} \left[ \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) [x_j^{(i)}] + \lambda \theta_j \right]$$

}



## **Logistic Regression Model**

## **Regularization-Cost Function**

# Regularization

*Repeat*

$$\theta_0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0) \quad for (j = 0)$$

$$\left. \begin{aligned} \theta_j &:= \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_j) \\ &\text{for } (j = 1, 2, 3 \dots, n) \end{aligned} \right\} \text{-----} (\textbf{Equation - A})$$

# Regularization

Calculation for  $\frac{\partial}{\partial \theta_0} J(\theta_0)$

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} (\log(h_{\theta}(x^{(i)}))) + (1 - y^{(i)}) \log(1 - (h_{\theta}(x^{(i)}))) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\frac{\partial}{\partial \theta_0} J(\theta_0)$$

$$= \frac{\partial}{\partial \theta_0} \left[ -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} (\log(h_{\theta}(x^{(i)}))) + (1 - y^{(i)}) \log(1 - (h_{\theta}(x^{(i)}))) \right] + \frac{\partial}{\partial \theta_0} \left[ \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2 \right] \right]$$

$$\frac{\partial}{\partial \theta_0} J(\theta_0) = \frac{1}{m} \left[ \sum_{i=1}^m \{y^{(i)} (1 - h_{\theta}(x^{(i)})) - (1 - y^{(i)}) (h_{\theta}(x^{(i)}))\} \right] x^{(i)}$$

# Regularization

**Calculation for  $\frac{\partial}{\partial \theta_j} J(\theta_j)$**

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} (\log(h_{\theta}(x^{(i)}))) + (1 - y^{(i)}) \log(1 - (h_{\theta}(x^{(i)}))) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta)$$

$$= \frac{\partial}{\partial \theta_j} \left[ -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} (\log(h_{\theta}(x^{(i)}))) + (1 - y^{(i)}) \log(1 - (h_{\theta}(x^{(i)}))) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2 \right]$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)} + \frac{\lambda \theta_j}{m}$$

# Regularization

- By putting the values in equation (A), we have

$$\begin{aligned}
 & \text{Repeat} \\
 & \{ \\
 \theta_0 &:= \theta_0 - \frac{\alpha}{m} \left[ \sum_{i=1}^m \{y^{(i)}(1 - h_{\theta}(x^{(i)}) - (1 - y^{(i)})(h_{\theta}(x^{(i)})))\} \right] x^{(i)} \\
 \theta_j &:= \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x^{(i)} + \frac{\lambda\theta_j}{m} \\
 & \}
 \end{aligned}$$



# Reference

---

- ▶ <https://www.coursera.org/>